

Name: Qiqian Fu
NetID: qiqianf2
Section: AL2

ECE 408/CS483 Milestone 3 Report

0. List Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images from your basic forward convolution kernel in milestone 2. This will act as your baseline this milestone. Note: **Do not** use batch size of 10k when you profile in `--queue rai_amd64_exclusive`. We have limited resources, so any tasks longer than 3 minutes will be killed. Your baseline M2 implementation should comfortably finish in 3 minutes with a batch size of 5k (About 1m35 seconds, with nv-nsight).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.17711ms	0.662171ms	0m5.974s	0.86
1000	1.68443ms	6.83812ms	0m54.411s	0.886
5000	8.33635ms	27.0691ms	4m24.008s	0.871

1. **Optimization 1: Using Streams to overlap computation with data transfer**

- a. Which optimization did you choose to implement and why did you choose that optimization technique.

The optimization I choose is to use stream because it allows concurrent execution, which can overlap data transfers between the host and device with kernel execution. Therefore, I think it can reduce the time of execution greatly.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

I think it should work and increase the performance greatly. As I mentioned, it can overlap data transfers between the host and device with kernel execution, which is quite time consuming in the baseline. It doesn't synergize with any of my previous optimizations.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	/	/	0m5.764s	0.86
1000	/	/	52.655s	0.886
5000	/	/	4m4.630s	0.871

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).

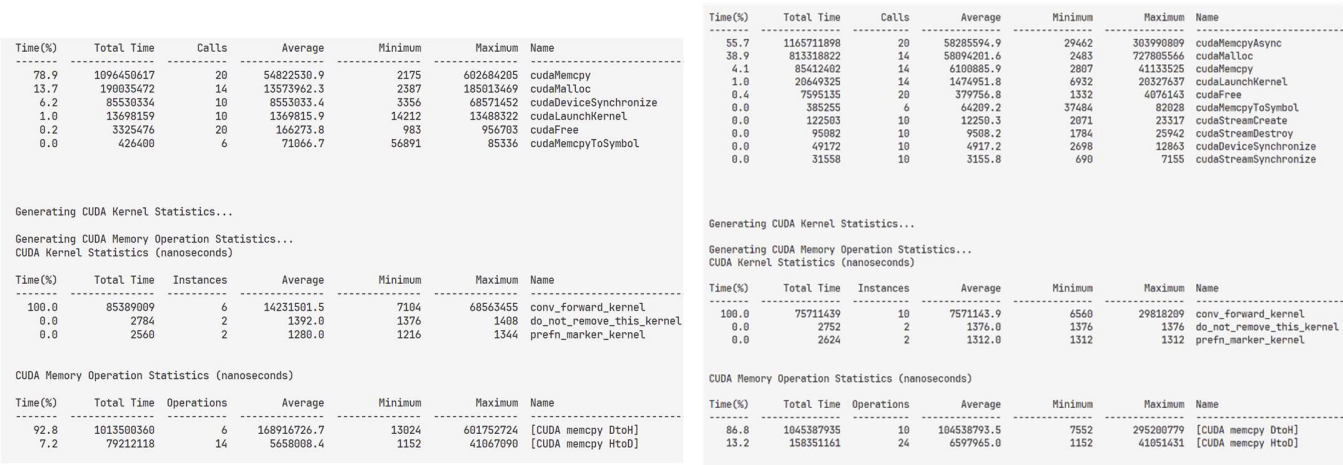
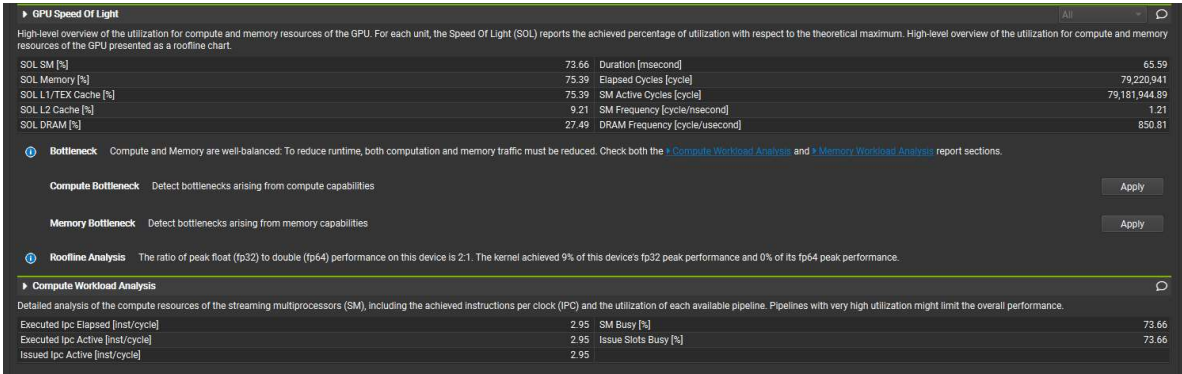
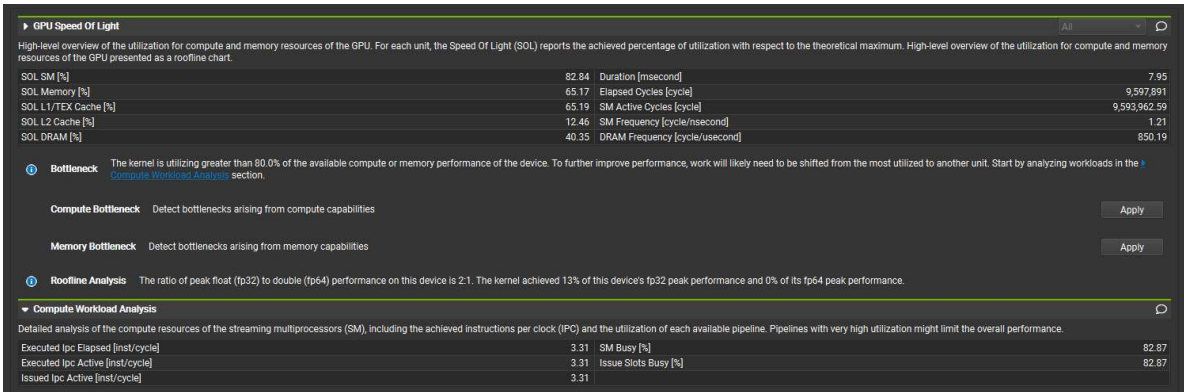


Fig 1

Yes, the implementation did improve the performance. The profiling result using *nsys* in Fig1 shows that the total time spending on 'conv_forward_kernel' does decrease a lot (from 85389009 to 75711439). Except for the operation time, I compared the profiling result using *Nsight-Compute*, and their differences have become even more pronounced which is shown in Fig2.



Baseline: Basic Convolution



Optimization: Stream
Fig 2

We can find that stream has approximately 9% higher streaming multiprocessor efficiency compared to baseline, indicating more effective use of the GPU's compute resources. Except for higher memory efficiency, stream has higher cache hit rates, better DRAM utilization, and increased IPC and Issue Slot utilization, indicating overall more effective memory handling and compute performance than baseline.

- e. What references did you use when implementing this technique?
 1. Nvidia official documentation: cuda-c-programming-guide chapter 11
 2. Nvidia slice named CUDA C/C++ Streams and Concurrency
 3. CSDN
 4. Stack overflow

2. Optimization 2: Tiled Shared Memory Convolution and Shared Memory Matrix Multiplication and Input Matrix Unrolling

- a. Which optimization did you choose to implement and why did you choose that optimization technique.

The second optimization I chose is Tiled Shared Memory Convolution and Shared Memory Matrix Multiplication and Input Matrix Unrolling. I think tiled shared memory convolution is a good way to reduce the times for global memory access and therefore decrease the bandwidth. However, it's not that easy to do tiled shared memory convolution in convolution compared with matrix multiplication because the tiled input size is normally larger than the mask size, not to mention that if there's stride, the size would be much larger. So the only solution is to unroll the input first then do the tiled shared memory convolution.

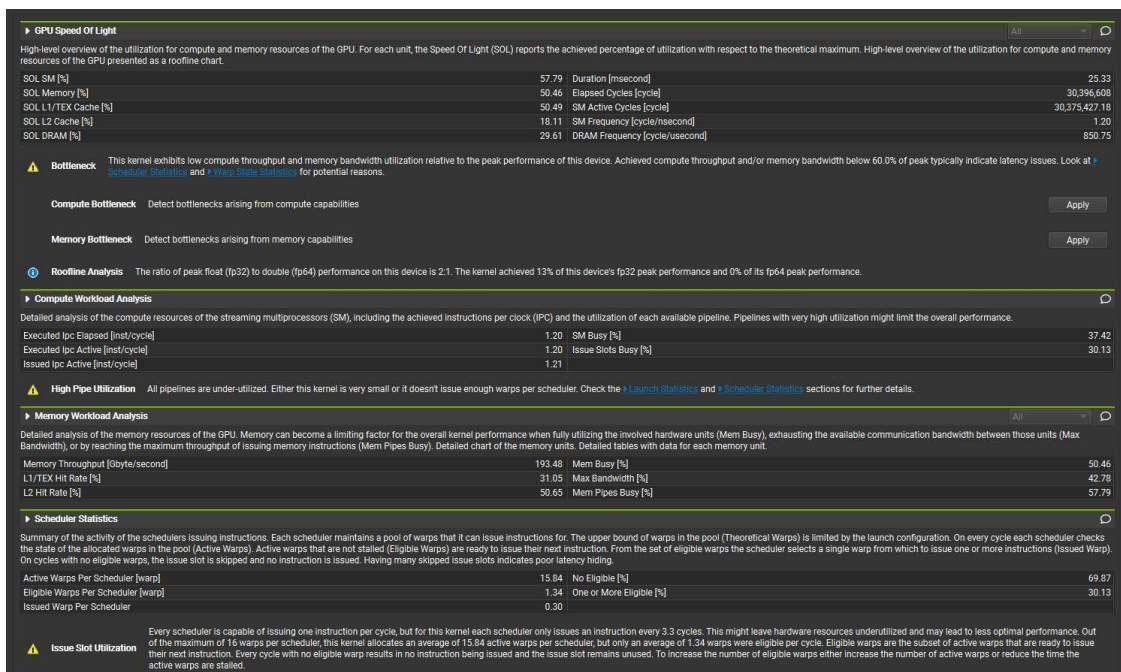
- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

Because it can reduce the times for global memory access and decrease the bandwidth. Additionally, it can make use of all the threads better. The optimization would definitely increase the performance of the forward convolution, because each thread will spend less time for data accessing. Logically, unroll synergizes with tiled shared memory, but I have represented them as a joint optimization because I believe they can only work in synergy, not independently.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.29814ms	0.42601ms	5.803s	0.86
1000	2.8649 ms	3.82974 ms	51.930s	0.886
5000	15.9825	22.5951	4m22.025s	0.871

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).



Optimization: Unrolling and Tiling

Fig 4

The optimization is not that good as my expected. Comparing the detailed information in Nsight-Computer Fig 4, we can find that the active warps per scheduler is much lower compared with the baseline in Fig 2. The main reason I think is that the if judgement in tiling occurs warp divergence. That's why the total execution time become even longer. But the duration does decrease, which means that the tiling does makes the calculation faster than baseline. And its memory throughput is also higher. That makes sense because the share memory of tiling can reduce the time for global memory access and thus help the kernel function to read the data faster than the baseline.

e. What references did you use when implementing this technique?

1. *CSDN*
2. *Stack overflow*

3. Optimization 3: Tuning with restrict and loop unrolling.

a. Which optimization did you choose to implement and why did you choose that optimization technique.

I choose to use tuning with restrict and loop unrolling because restrict can be used as a keyword to tell the compiler that the no pointer will point to the same place thus allowing the compiler not to worry data overlap or pointer aliasing. What's more, loop unrolling can help increase the efficiency of each thread's work since the mask size is fixed in this project.

b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

I think it will help the thread working more quickly and increase the performance of the forward convolution. No, it's a separate optimization and if it works, it should increase the performance of the base-line implementation.

c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.2761ms	0.3929ms	5.603s	0.86
1000	1.3397 ms	5.6390 ms	0m52.411s	0.886
5000	6.6417 ms	27.0089 ms	4m11.939s	0.871

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).

Yes, it was successful.

time	seconds	seconds	calls	s/call	s/call	name
10.20	3.45	3.45	2	1.73	14.21	MaxPooling::forward(Eigen::Matrix<float, -1, -1, 0, -1, -1> const&)
Baseline						
time	seconds	seconds	calls	s/call	s/call	name
9.44	3.05	3.05	2	1.53	13.47	MaxPooling::forward(Eigen::Matrix<float, -1, -1, 0, -1, -1> const&)
Tuning with restrict and loop unrolling						

Comparing the time spending on forward with nsys, we can see that it takes less time using restrict and unrolling.

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name	Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
78.5	1131322890	20	56566144.5	27860	607804322	cudaMemcpy	78.9	1052231683	20	52611584.1	2400	572637448	cudaMemcpy
13.7	197561570	20	9878078.5	2604	193495867	cudaMalloc	14.3	190122728	14	13580194.9	2520	184961748	cudaMalloc
6.5	93063861	10	9306386.1	3571	74861030	cudaDeviceSynchronize	5.3	70945259	10	7094525.9	5366	57326542	cudaDeviceSynchronize
1.1	15916403	10	1591640.3	24610	15662091	cudaLaunchKernel	1.2	16026232	10	1602623.2	14816	15809887	cudaLaunchKernel
0.3	3725611	20	186280.5	6380	1057003	cudaFree	0.2	3150415	20	157520.8	1014	970342	cudaFree
							0.0	423372	6	70562.0	55141	85260	cudaMemcpyToSymbol

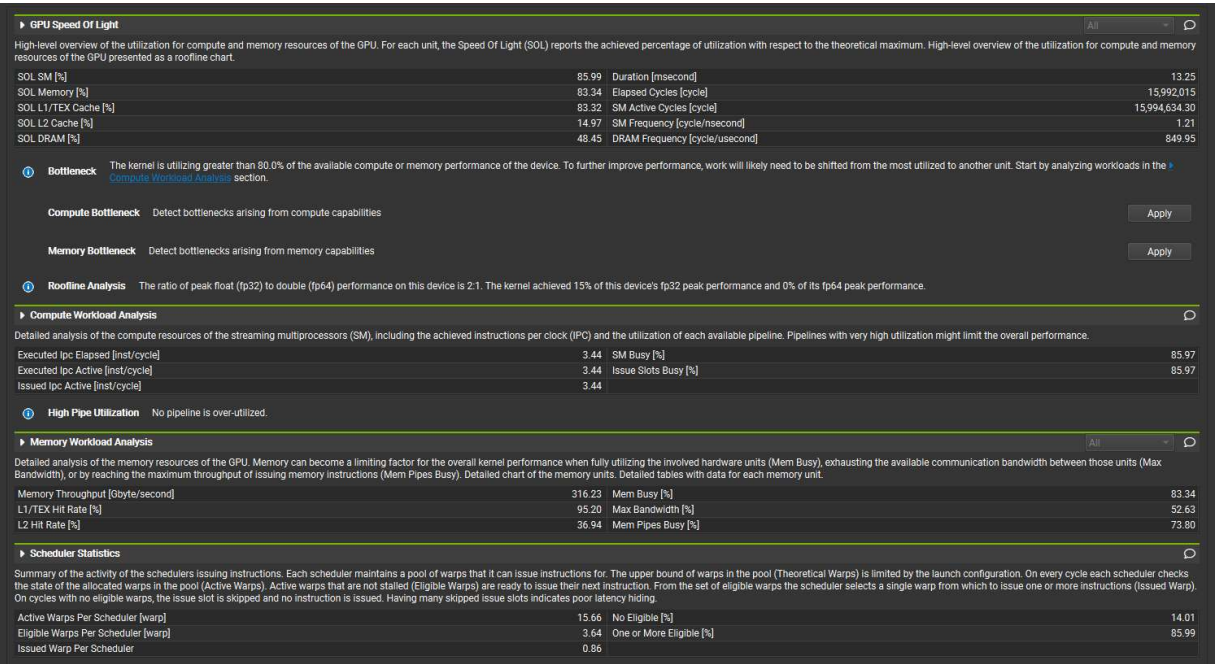
Generating CUDA Kernel Statistics...							Generating CUDA Kernel Statistics...						
Generating CUDA Memory Operation Statistics...							Generating CUDA Memory Operation Statistics...						
CUDA Kernel Statistics (nanoseconds)							CUDA Kernel Statistics (nanoseconds)						
Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name	Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	93038595	6	15506432.5	7392	74853254	conv_forward_kernel	100.0	70786930	6	11797821.7	7168	57314626	conv_forward_kernel
0.0	2784	2	1392.0	1376	1408	do_not_remove_this_kernel	0.0	2816	2	1408.0	1408	1408	do_not_remove_this_kernel
0.0	2656	2	1328.0	1312	1344	prefn_marker_kernel	0.0	2624	2	1312.0	1280	1344	prefn_marker_kernel

Baseline: Basic Convolution

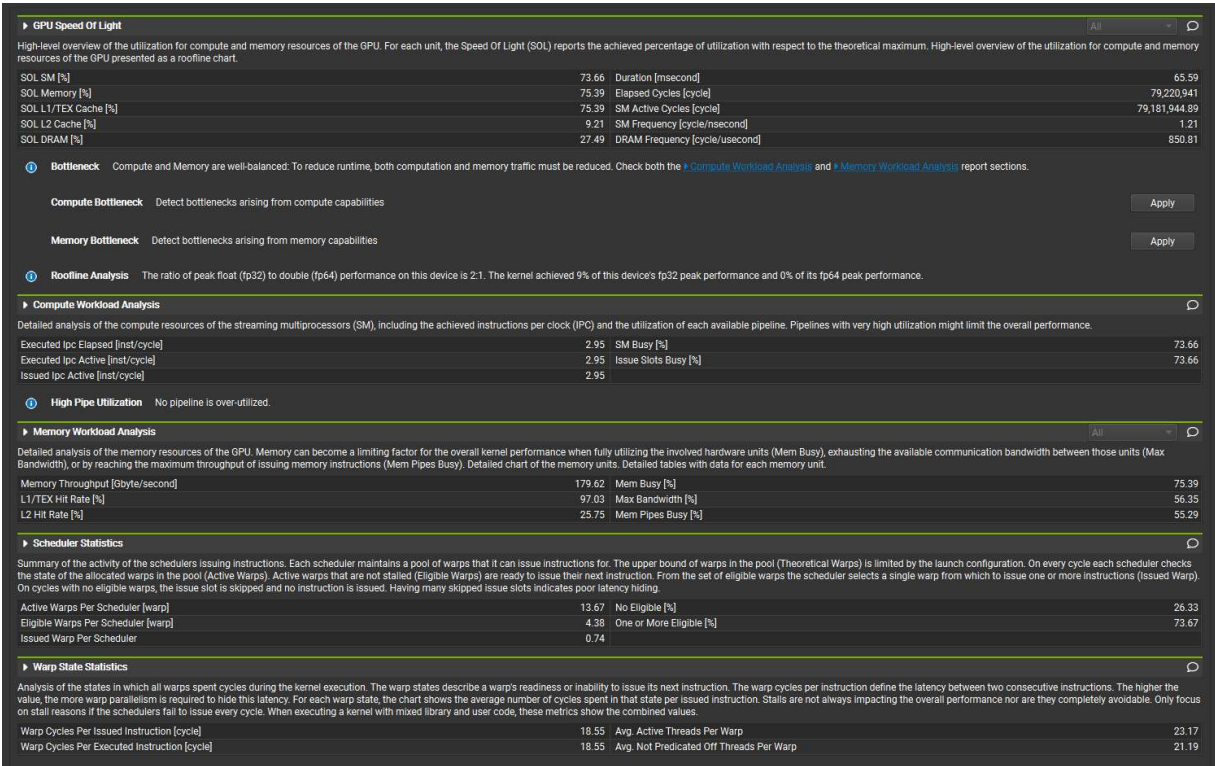
Optimization: Tuning and Restrict

Fig 4

And in Fig 4 we can see that the total time spending on conv_forward_kernel is decreased tremendously. In Fig 5, we show the detailed information of the optimization using Nsight-Computer. Comparing with the Basic Convolution in Fig 2, we can find that the SOL SM increase from 74% to 86%, and memory throughput increase from 179% to 316%. I think SOL SM shows that the tuning does help streaming multiprocessor to effectively use of the GPU's compute resources, and memory throughput means the rate at which data is being read from or written to memory by the GPU, which demonstrates that the reading data speed is much higher with the restrict keyword.



Optimization: Tuning and Restrict



Baseline: Basic Convolution

Fig 5

e. What references did you use when implementing this technique?

1. CSDN
2. Open Liberty
3. Stack Overflow

4. Optimization 4: Weight matrix (kernel values) in constant memory

a. Which optimization did you choose to implement and why did you choose that optimization technique.

I chose to use weight matrix (kernel values) in constant memory because constant memory access takes about 5 times longer than local memory access which is much less than 20 times for global memory access.

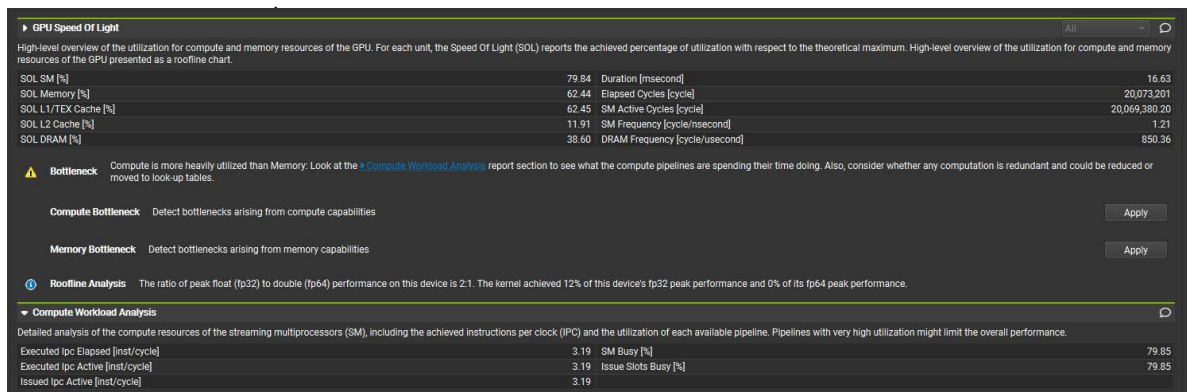
b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

I think it will reduce the operation time because like I said, constant memory access is faster than global memory access, especially for mask data which will be read many times by threads. No, it doesn't synergize with any of your previous optimizations.

c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.1859ms	0.6201ms	0m5.848s	0.86
1000	1.5341ms	6.6709ms	0m53.190s	0.886
5000	8.0699ms	25.0621ms	4m20.179s	0.871

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).



Optimization: Constant Memory

Fig 6

Compared with the baseline in Fig 2 using *Nsight-Compute*, the SOL SM, SOL Memory, SOL DRAM increases slightly. That's due to the constant memory. And the SM busy, Executed lpc Elapsed are the same, which means constant memory doesn't help reduce the computation workload, it only helps each thread reads data faster.

- e. What references did you use when implementing this technique?
1. *Textbook Programming Massively Parallel Processors A Hands-On Approach* (David B. Kirk, Wen-Mei W Hwu) (Z-Library).pdf
 2. MP5

5. Optimization 5: Sweeping various parameters to find best values

- a. Which optimization did you choose to implement and why did you choose that optimization technique.

I chose Sweeping various parameters to find best values because I think good parameters helps the streaming multiprocessor to be more rational in planning of computation resources

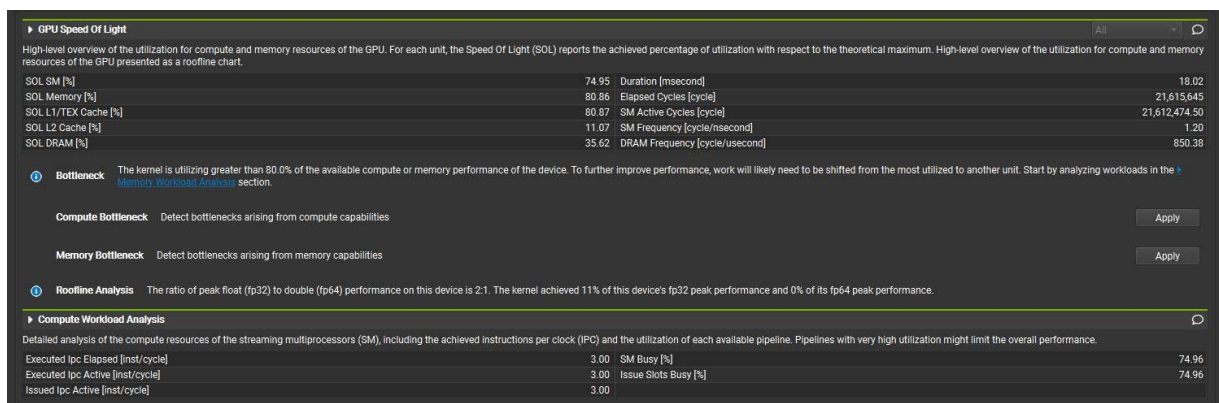
- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

I think it will have some help because if the block size is 32 or times of 32, then the SM can better make use of the computation resources. No, it doesn't synergize with any of your previous optimizations.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.1894ms	0.6452ms	0m5.990s	0.86
1000	1.5418ms	6.7988ms	0m55.891s	0.886
5000	8.7167ms	24.9920ms	4m27.653s	0.871

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).



Optimization: Sweeping Various

Fig 7

The implementation is successful in improving performance. As I expected, setting blockdim to 32 helps the model to increase executed lpc Elapsed and SM Busy. Because it avoid warp divergence, makes threads to run concurrently.

e. What references did you use when implementing this technique?

1. *Textbook Programming Massively Parallel Processors A Hands-On Approach*
(David B. Kirk, Wen-Mei W Hwu) (Z-Library).pdf
2. MP2